

Bayesian MAP Model Selection of Chain Event Graphs

G. Freeman^a, J.Q. Smith^a

^a*Department of Statistics, University of Warwick, Coventry, CV4 7AL*

Abstract

The class of chain event graph models is a generalisation of the class of discrete Bayesian networks, retaining most of the structural advantages of the Bayesian network for model interrogation, propagation and learning, while more naturally encoding asymmetric state spaces and the order in which events happen. In this paper we demonstrate how with complete sampling, conjugate closed form model selection based on product Dirichlet priors is possible, and prove that suitable homogeneity assumptions characterise the product Dirichlet prior on this class of models. We demonstrate our techniques using two educational examples.

Key words: chain event graphs, Bayesian model selection, Dirichlet distribution

1. Introduction

Bayesian networks (BNs) are currently one of the most widely used graphical models for representing and analysing finite discrete graphical multivariate distributions with their explicit coding of conditional independence relationships between a system's variables [1, 2]. However, despite their power and usefulness, it has long been known that BNs cannot fully or efficiently represent certain common scenarios. These include situations where the state space of a variable is known to depend on other variables, or where the conditional independence between variables is itself dependent on the values of other variables. Some examples of such latter scenarios are given by Poole and Zhang [3]. In order to overcome such deficiencies, enhancements have been proposed to the basic Bayesian network in order to create so-called “context-specific” Bayesian networks [3]. These have their own problems, however: either they represent too much of the information about a model in a non-graphical way, thus undermining the rationale for using a graphical model in the first place, or they struggle to represent a general class of models efficiently. Other graphical approaches that seek to account for “context-specific” beliefs suffer from similar problems.

This has led to the proposal of a new graphical model — the chain event graph (CEGs) — which first propounded in [4]. As well as solving the aforementioned problems associated with Bayesian networks and related graphical models, CEGs are able, not unrelatedly, to encode far more efficiently the common structure in which models are elicited — as asymmetric processes — in a single graph. To this end, CEGs are based not on Bayesian networks, but on event trees (ETs) [5]. Event trees are trees where nodes represent situations — i.e. scenarios in which a unit might find itself — and each node's extending edges represent possible future situations that can develop

Email addresses: g.freeman@warwick.ac.uk (G. Freeman), j.q.smith@warwick.ac.uk (J.Q. Smith)

from the current one. It follows that every atom of the event space is encoded by exactly one root-to-leaf path, and each root-to-leaf path corresponds to exactly one atomic event. It has been argued that ETs are expressive frameworks to directly and accurately represent beliefs about a process, particularly when the model is described most naturally, as in the example below, through how situations might unfold [5]. However, as explained in [4], ETs can contain excessive redundancy in their structure, with subtrees describing probabilistically isomorphic unfoldings of situations being represented separately. They are also unable to explicitly express a model's non-trivial conditional independences. The CEG deals with these shortcomings by combining the subtrees that describe identical subprocesses (see [4] for further details), so that the CEG derived from a particular ET has a simpler topology while in turn expressing more conditional independence statements than is possible through an ET.

We illustrate the construction and the types of symmetries it is possible to code using a CEG with the following running example.

Example 1. *Successful students on a one year programme study components A and B , but not everyone will study the components in the same order: each student will be allocated to study either module A or B for the first 6 months and then the other component for the final 6 months. After the first 6 months each student will be examined on their allocated module and be awarded a distinction (denoted with D), a pass (P) or a fail (F), with an automatic opportunity to resit the module in the last case. If they resit then they can pass and be allowed to proceed to the other component of their course, or fail again and be permanently withdrawn from the programme. Students who have succeeded in proceeding to the second module can again either fail, pass or be awarded a distinction. On this second round, however, there is no possibility of resitting if the component is failed. With an obvious extension of the labelling, this system can be depicted by the event tree given in Figure 1.*

To specify a full probability distribution for this model it is sufficient to only specify the distributions associated with the unfolding of each situation a student might reach. However, in many applications it is often natural to hypothesise a model where the distribution associated with the unfolding from one situation is assumed identical to another. Situations that are thus hypothesised to have the same transition probabilities to their children are said to be in the same *stage*. Thus in Example 1 suppose that as well as subscribing to the ET of Figure 1 we want to consider a model also embodying the following three hypotheses:

1. The chances of doing well in the second component are the same whether the student passed first time or after a resit.
2. The components A and B are equally hard.
3. The distribution of marks for the second component is unaffected by whether students passed or got a distinction for the first component.

These hypotheses can be identified with a partitioning of the non-leaf nodes (situations). In Figure 1 the set of situations is

$$\mathcal{S} = \{V_0, A, B, P_{1,A}, P_{1,B}, D_{1,A}, D_{1,B}, F_{1,A}, F_{1,B}, P_{R,A}, P_{R,B}\}.$$

The partition C of \mathcal{S} that encodes exactly the above three hypotheses consists of the stages $u_1 = \{A, B\}$, $u_2 = \{F_{1,A}, F_{1,B}\}$, and $u_3 = \{P_{1,A}, P_{1,B}, P_{R,A}, P_{R,B}, D_{1,A}, D_{1,B}\}$ together with the singleton $u_0 = \{V_0\}$. Thus the second stage u_2 , for example, implies that the probabilities

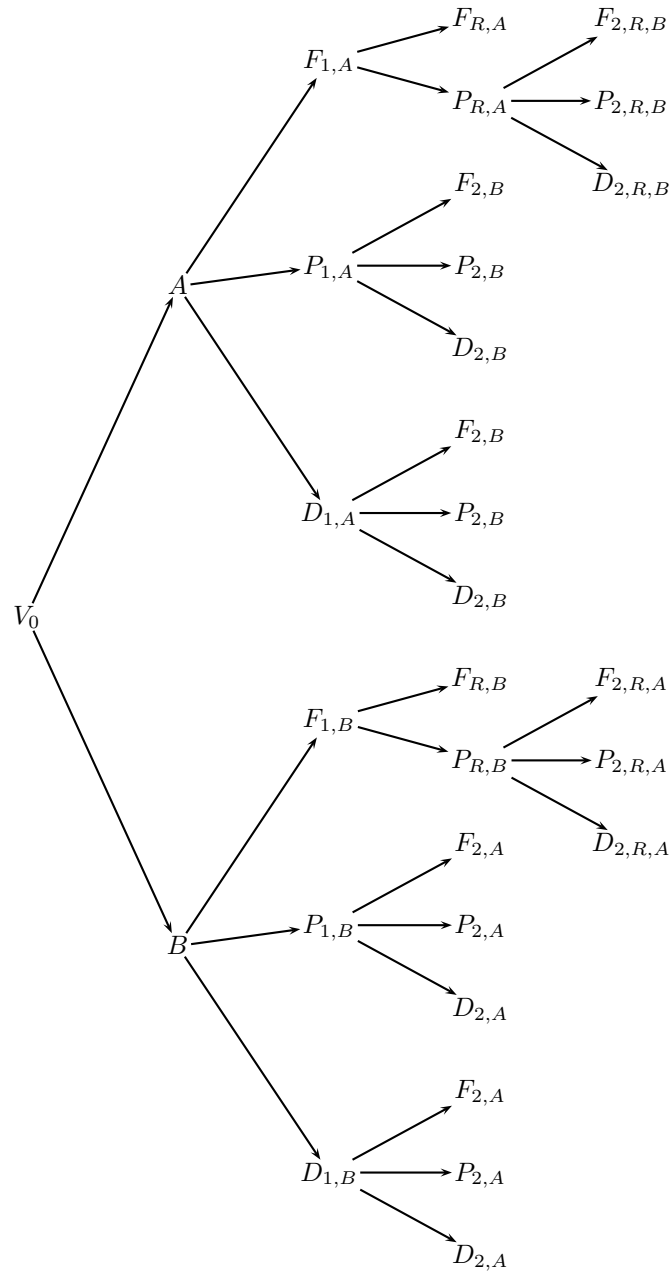


Figure 1: Event tree of a student's potential progress through a hypothetical course described in Example 1. Each non-leaf node represents a juncture at which a random event will take place, with the selection of possible outcomes represented by the edges emanating from that node. Each edge distribution is defined conditional on the path passed through earlier in the tree to reach the specific node.

on the edges $(F_{1,B}, F_{R,B})$ and $(F_{1,A}, F_{R,A})$ are equal, as are the probabilities on $(F_{1,B}, P_{R,B})$ and $(F_{1,A}, P_{R,A})$. Clearly the joint probability distribution of the model – whose atoms are the root to leaf paths of the tree – is determined by the conditional probabilities associated with the stages. A CEG is the graph that is constructed to encode a model that can be specified through an event tree combined with a partitioning of its situations into stages.

In this paper we suppose that we are in a context similar to that of Example 1, where, for any possible model, the sample space of the problem must be consistent with a single event tree, but where on the basis of a sample of students' records we want to select one of a number of different possible CEG models, i.e. we want to find the “best” partitioning of the situations into stages. We take a Bayesian approach to this problem and choose the model with the highest posterior probability — the Maximum A Posteriori (MAP) model. This is the simplest and possibly most common Bayesian model selection method, advocated by, for example, Dennison et al [6], Castelo [7], and Heckerman [8], the latter two specifically for Bayesian network selection.

The paper is structured as follows. In the next section we review the definitions of event trees and CEGs. In Section 3 we develop the theory of how conjugate learning of CEGs is performed. In Section 4 we apply this theory by using the posterior probability of a CEG as its score in a model search algorithm that is derived using an analogous procedure to the model selection of BNs. We characterise the product Dirichlet distribution as a prior distribution for the CEGs' parameters under particular homogeneity conditions. In Section 5 the algorithm is used to discover a good explanatory model for real students' exam results. We finish with a discussion.

2. Definitions of event trees and chain event graphs

In this section we briefly define the event tree and chain event graph. We refer the interested reader to [4] for further discussion and more detail concerning their construction. Bayesian networks, which will be referenced throughout the paper, have been defined many times before. See [8] for an overview.

2.1. Event Trees

Let $T = (V(T), E(T))$ be a directed tree where $V(T)$ is its node set and $E(T)$ its edge set. Let $S(T) = \{v : v \in V(T) - L(T)\}$ be the set of SITUATIONS of T , where $L(T)$ is the set of LEAF (or TERMINAL) nodes. Furthermore, define $\mathbb{X} = \{\lambda(v_0, v) : v \in V(T) \setminus S(T)\}$, where $\lambda(a, b)$ is the path from node a to node b , and v_0 is the root node, so that \mathbb{X} is the set of root-to-leaf paths of T . Each element of \mathbb{X} is called an ATOMIC EVENT, each one corresponding to a possible unfolding of events through time by using the partial ordering induced by the paths. Let $\mathbb{X}(v)$ denote the set of children of $v \in V(T)$. In an event tree, each situation $v \in S(T)$ has an associated random variable $X(v)$ with sample space $\mathbb{X}(v)$, defined conditional on having reached v . The distribution of $X(v)$ is determined by the PRIMITIVE PROBABILITIES $\{\pi(v'|v) = p(X(v) = v') : v' \in \mathbb{X}(v)\}$. With random variables on the same path being mutually independent, the joint probability of events on a path can be calculated by multiplying the appropriate primitive probabilities together. Each primitive probability $\pi(v'|v)$ is a colour for the directed edge $e = (v, v')$, so that we can have $\pi(e) = \pi(v'|v)$.

Example 2. Figure 2 shows a tree for two Bernoulli random variables, X and Y , with X occurring before Y . In an educational example X could be the indicator variable of a student passing one module, and Y the indicator variable for a subsequent module.




Figure 2: Simple event tree. The non-zero-probability events in the joint probability distribution of two Bernoulli random variables, X and Y , with X observed before Y , can be represented by this tree. Here, all four joint states are possible, because there are four root-to-leaf paths through the nodes.




Figure 3: Floret of v . This subtree represents both the random variable $X(v)$ and its state space $\mathbb{X}(v)$.

Here we have random variables $X(v_0) = X$, $X(v_1) = Y|(X = 0)$ and $X(v_2) = Y|(X = 1)$, and primitive probabilities $\pi(v_1|v_0) = p(X = 0)$, $\pi(v_3|v_1) = p(Y = 0|X = 0)$ and so on for every other edge. Joint probabilities can be found by multiplying primitive probabilities along a path, e.g. $p(X = 0, Y = 0) = p(X = 0)p(Y = 0|X = 0) = \pi(v_1|v_0)\pi(v_3|v_1)$ as v_0 and v_1 are on a path.

2.2. Chain Event Graphs

Starting with an event tree T , define a FLORET of $v \in S(T)$ as

$$\mathcal{F}(v, T) = (V(\mathcal{F}(v, T)), E(\mathcal{F}(v, T)))$$

where $V(\mathcal{F}(v, T)) = \{v\} \cup \{v' \in V(T) : (v, v') \in E(T)\}$ and $E(\mathcal{F}(v, T)) = \{e \in E(T) : e = (v, v')\}$. The floret of a vertex v is thus a sub-tree consisting of v , its children, and the edges connecting v and its children, as shown in Figure 3. This represents, as defined in section 2.1, the random variable $X(v)$ and its sample space $\mathbb{X}(v)$.

One of the redundancies that can be eliminated from an ET is that of the florets' edges of two situations, v and v' say, which have identical associated edge probabilities despite being defined by different conditioning paths. We say these two situations are at the same STAGE. This concept is formally defined as follows.

Definition 3. Two situations $v, v' \in S(T)$ are in the same stage u if and only if $X(v)$ and $X(v')$ have the same distribution under a bijection

$$\psi_u(v, v') : E(\mathcal{F}(v, T)) \rightarrow E(\mathcal{F}(v', T))$$

i.e.

$$\psi_u(v, v') : \mathbb{X}(v) \rightarrow \mathbb{X}(v')$$

The set of stages of an ET T is written $J(T)$. This set partitions the set of situations $S(T)$.

We can construct a STAGED TREE $\mathcal{G}(T, L(T))$ with $V(\mathcal{G}) = V(T)$, $E(\mathcal{G}) = E(T)$, and colour its edges such that:

- If $v \in u$ and u contains no other vertices, then all $(v, v^*) \in E(\mathcal{G})$ are left uncoloured;
- If $v \in u$ and u contains other vertices, then all $(v, v^*) \in E(\mathcal{G})$ are coloured; and
- Whenever $e(v, v^*) \mapsto e(v', v'^*)$ under $\psi_u(v, v')$, then the two edges must have the same colour.

There is another type of situation that is of further interest. When the whole development from two situations v and v' have identical distributions, i.e. there exists a bijection between their respective subtrees similar to that between stages as defined in Definition 2.2, then the situations are said to be in the same POSITION. This is defined formally as follows.

Definition 4. *Two situations $v, v' \in S(T)$ are in the same position w if and only if there exists a bijection*

$$\phi_w(v, v') : \Lambda(v, T) \rightarrow \Lambda(v', T)$$

where $\Lambda(v, T)$ is the set of paths in T from v to a leaf node of T , such that

- all edges in all of the paths in $\Lambda(v, T)$ and $\Lambda(v', T)$ are coloured in $\mathcal{G}(T, L(T))$; and
- for every path $\lambda(v) \in \Lambda(v, T)$, the ordered sequence of colours in $\lambda(v)$ equals the ordered sequence of colours in $\lambda(v') := \phi_w(v, T)(\lambda(v)) \in \Lambda(v', T)$

This ensures that when v and v' are in the same position, then under the map $\phi_w(v, v')$ future development from either node follows identical probability distributions.

We denote the set of positions as $K(T)$. Positions are an obvious way of equating situations, because the different conditioning variables of different nodes in the same position have no effect on any subsequent development. It is clear that $K(T)$ is a finer partition of $V(T)$ than $J(T)$, and indeed that $J(T)$ partitions $K(T)$, as situations in the same position will also be in the same stage.

We now use stages and positions to compress the event tree into a chain event graph. First, the PROBABILITY GRAPH of the event tree

$$\mathcal{H}(\mathcal{G}(T)) = \mathcal{H}(T) = (V(\mathcal{H}), E(\mathcal{H}))$$

is drawn, where $V(\mathcal{H}) = K(T) \cup \{w_\infty\}$ and $E(\mathcal{H})$ is constructed as follows.

- For each pair of positions $w, w' \in K(T)$, if there exists $v, v' \in S(T)$ such that $v \in w, v' \in w'$ and $e(v, v') \in E(T)$, then an associated edge $e(w, w') \in E(\mathcal{H})$ is drawn. Furthermore, if for a position w there exists $v \in S(T)$, $v' \in L(T)$ and $e(v, v') \in E(T)$ such that $v \in w$, then an associated edge $e(w, w_\infty) \in E(\mathcal{H})$ is drawn.
- The colour of this edge, $e(w, w')$, is the same as the colour of the associated edge $e(v, v')$.

Now the CEG can finally be constructed by taking the probability graph $\mathcal{H}(T)$ and connecting the positions that are in the same stage using undirected edges: Let $\mathcal{C}(T)$ be a mixed graph with vertex set $V(\mathcal{C}) = V(\mathcal{H})$, directed edge set $E_d(\mathcal{C}) = E(\mathcal{H})$, and undirected edge set $E_u(\mathcal{C}) = \{(w, w') : u(w) = u(w'), w, w' \in V(\mathcal{C})\}$.

An example of a CEG that could be constructed from the event tree in Figure 1 is shown in Figure 5.1.

3. Conjugate learning of CEGs

One convenient property of CEGs is that conjugate updating of the model parameters proceeds in a closely analogous fashion to that on a BN. Conjugacy is a crucial part of the model selection algorithm that will be described in Section 4, because it leads to closed form expressions for the posterior probabilities of candidate CEGs. This in turn makes it possible to search the often very large model space quickly to find optimal models. We demonstrate here how a conjugate analysis on a CEG proceeds.

Let a CEG C have set of stages $J(C) = \{u_1, \dots, u_k\}$, and let each stage u_i have k_i emanating edges (labelled e_1, \dots, e_{k_i}) with associated probability vector $\boldsymbol{\pi}_i = (\pi_{i1}, \pi_{i2}, \dots, \pi_{ik_i})'$ (where $\sum_{j=1}^{k_i} \pi_{ij} = 1$ and $\pi_{ij} > 0$ for $j \in \{1, \dots, k_i\}$). Then, under random sampling, the likelihood of the CEG can be decomposed into a product of the likelihood of each probability vector, i.e.

$$p(\mathbf{x}|\boldsymbol{\pi}, C) = \prod_{i=1}^k p_i(\mathbf{x}_i|\boldsymbol{\pi}_i, C)$$

where $\boldsymbol{\pi} = \{\boldsymbol{\pi}_1, \boldsymbol{\pi}_2, \dots, \boldsymbol{\pi}_k\}$, and $\mathbf{x} = \{\mathbf{x}_1, \dots, \mathbf{x}_k\}$ is the complete sample data such that each $\mathbf{x}_i = (x_{i1}, \dots, x_{ik_i})'$ is the vector of the number of units in the sample (for example, the students in Example 1) that start in stage u_i and move to the stage at the end of edge e_{ij} for $j \in \{1, \dots, k_i\}$.

If it is further assumed that $\mathbf{x}_i \perp \mathbf{x}_j | \boldsymbol{\pi}, \forall i \neq j$ then

$$p_i(\mathbf{x}_i|\boldsymbol{\pi}_i, C) = \prod_{j=1}^{k_i} \pi_{ij}^{x_{ij}} \quad (1)$$

Thus, just as for the analogous situation with BNs, the likelihood of a random sample also separates over the components of $\boldsymbol{\pi}$. With BNs, a common modelling assumption is of local and global independence of the probability parameters [9]; the corresponding assumption here is that the parameters $\boldsymbol{\pi}_1, \boldsymbol{\pi}_2, \dots, \boldsymbol{\pi}_k$ of $\boldsymbol{\pi}$ are all mutually independent a priori. It will then follow, with the separable likelihood, that they will also be independent a posteriori.

If the probabilities $\boldsymbol{\pi}_i$ are assigned a Dirichlet distribution, $\text{Dir}(\boldsymbol{\alpha}_i)$, a priori, where $\boldsymbol{\alpha}_i = (\alpha_{i1}, \alpha_{i2}, \dots, \alpha_{ik_i})'$, so that for values of π_{ij} such that $\sum_{j=1}^{k_i} \pi_{ij} = 1$ and $\pi_{ij} > 0$ for $1 \leq j \leq k_i$, the density of $\boldsymbol{\pi}_i$, $q_i(\boldsymbol{\pi}_i|C)$, can be written

$$q_i(\boldsymbol{\pi}_i|C) = \frac{\Gamma(\alpha_{i1} + \dots + \alpha_{ik_i})}{\Gamma(\alpha_{i1}) \dots \Gamma(\alpha_{ik_i})} \prod_{j=1}^{k_i} \pi_{ij}^{\alpha_{ij}-1}$$

where $\Gamma(z) = \int_0^\infty t^{z-1} e^{-t} dt$ is the Gamma function. It then follows that $\boldsymbol{\pi}_i | \mathbf{x}$ ($= \boldsymbol{\pi}_i | \mathbf{x}_i$) also has a Dirichlet distribution, $\text{Dir}(\boldsymbol{\alpha}_i^*)$, a posteriori, where $\boldsymbol{\alpha}_i^* = (\alpha_{i1}^*, \dots, \alpha_{ik_i}^*)'$, $\alpha_{ij}^* = \alpha_{ij} + x_{ij}$ for $1 \leq j \leq k_i, 1 \leq i \leq k$. The marginal likelihood of this model can be written down explicitly as the function of the prior and posterior Dirichlet parameters:

$$p(\mathbf{x}|C) = \prod_{i=1}^k \left[\frac{\Gamma(\sum_j \alpha_{ij})}{\Gamma(\sum_j \alpha_{ij}^*)} \prod_{j=1}^{k_i} \frac{\Gamma(\alpha_{ij}^*)}{\Gamma(\alpha_{ij})} \right].$$

The computationally more useful logarithm of the marginal likelihood is therefore a linear combination of functions of α_{ij} and α_{ij}^* . Explicitly,

$$\log p(\mathbf{x}|C) = \sum_{i=1}^k [s(\alpha_i) - s(\alpha_i^*)] + \sum_{i=1}^k [t(\alpha_i^*) - t(\alpha_i)] \quad (2)$$

where for any vector $\mathbf{c} = (c_1, c_2, \dots, c_n)'$,

$$s(\mathbf{c}) = \log \Gamma\left(\sum_{v=1}^n c_v\right) \text{ and } t(\mathbf{c}) = \sum_{v=1}^n \log \Gamma(c_v) \quad (3)$$

So the posterior probability of a CEG C after observing \mathbf{x} , $q(C|\mathbf{x})$, can be calculated using Bayes' Theorem, given a prior probability $q(C)$:

$$\log q(C|\mathbf{x}) = \log p(\mathbf{x}|C) + \log q(C) + K \quad (4)$$

for some value K which does not depend on C . This is the SCORE that will be used when searching over the candidate set of CEGs for the model that best describes the data.

4. A Local Search Algorithm for Chain Event Graphs

4.1. Preliminaries

With the log marginal posterior probability of a CEG model, $\log q(C|\mathbf{x})$, as its score, searching for the highest-scoring CEG in the set of all candidate models is equivalent to trying to find the Maximum A Posteriori (MAP) model [10]. The intuitive approach for searching \mathbf{C} , the candidate set of CEGs — calculating $q(C|\mathbf{x})$ (or $\log q(C|\mathbf{x})$) for every $C \in \mathbf{C}$ and choosing $C^* := \max_C q(C|\mathbf{x}) = \max_C \log q(C|\mathbf{x})$ — is infeasible for any but the most trivial problems. We describe in this section an algorithm for efficiently searching the model space by reformulating the model search problem as a clustering problem.

As mentioned in Section 2.2, every CEG that can be formed from a given event tree can be identified exactly with a partition of the event tree's nodes into stages. The coarsest partition C_∞ has all nodes with k outgoing edges in the same stage, u_k ; the finest partition C_0 has each situation in its own stage, except for the trivial cases of those nodes with only one outgoing edge. Defined this way, the search for the highest-scoring CEG is equivalent to searching for the highest-scoring clustering of stages.

Various Bayesian clustering algorithms exist [11], including many involving MCMC [12]. We show here how to implement an Bayesian agglomerative hierarchical clustering (AHC) exact algorithm related to that of Heard et al [13]. The AHC algorithm here is a local search algorithm that begins with the finest partition of the nodes of the underlying ET model (called C_0 above and henceforth) and seeks at each step to find the two nodes that will yield the highest-scoring CEG if combined.

Some optional steps can be taken to simplify the search, which we will implement here. The first of these involves the calculation of the scores of the proposed models in the algorithm. By assuming that the probability distributions of stages that are formed from the same nodes of the underlying ET are equal in all CEGs, i.e. $p(\mathbf{x}_i|\boldsymbol{\pi}_i, C_1) = p(\mathbf{x}_i|\boldsymbol{\pi}_i, C_2), \forall C_1, C_2 \in \mathbf{C}$, it becomes more efficient to calculate the differences of model scores, i.e. the logarithms of the relevant Bayes factors, than to calculate the two individual model scores absolutely. This is because, if for two

CEGs their stage sets $J(C_1)$ and $J(C_2)$ differ only in that stages $u_{1a}, u_{1b} \in C_1$ are combined into $u_{2c} \in C_2$, with all other stages unchanged, then the calculation of the logarithm of their posterior Bayes factor depends only on the stages involved; using the notation of Equation (3),

$$\log \frac{q(C_1|\mathbf{x})}{q(C_2|\mathbf{x})} = \log q(C_1|\mathbf{x}) - \log q(C_2|\mathbf{x}) \quad (5)$$

$$= \log q(C_1) - \log q(C_2) + \log q(\mathbf{x}|C_1) - \log q(\mathbf{x}|C_2) \quad (6)$$

$$= \log q(C_1) - \log q(C_2) + \sum_i [s(\alpha_{1i}) - s(\alpha_{1i}^*)] + \sum_i [t(\alpha_{1i}^*) - t(\alpha_{1i})] \\ - \sum_i [s(\alpha_{2i}) - s(\alpha_{2i}^*)] - \sum_i [t(\alpha_{2i}^*) - t(\alpha_{2i})] \quad (7)$$

$$= \log q(C_1) - \log q(C_2) + s(\alpha_{1a}) - s(\alpha_{1a}^*) + t(\alpha_{1a}^*) - t(\alpha_{1a}) \\ + s(\alpha_{1b}) - s(\alpha_{1b}^*) + t(\alpha_{1b}^*) - t(\alpha_{1b}) \\ - s(\alpha_{2c}) + s(\alpha_{2c}^*) - t(\alpha_{2c}^*) + t(\alpha_{2c}) \quad (8)$$

Using the trivial result that for any three CEGs

$$\log q(C_3|\mathbf{x}) - \log q(C_2|\mathbf{x}) = [\log q(C_3|\mathbf{x}) - \log q(C_1|\mathbf{x})] - [\log q(C_2|\mathbf{x}) - \log q(C_1|\mathbf{x})],$$

it can be seen that in the course of the AHC algorithm, comparing two proposal CEGs from the current CEG can be done equivalently by comparing their log Bayes factors with the current CEG, which as shown above requires fewer calculations.

The calculation of the score for each CEG C , as shown by Equation (4), shows that it is formed of two components: the prior probability of the CEG being the true model and the marginal likelihood of the data. These must therefore be set before the algorithm can be run, and it is here that the other simplifications are made.

4.2. The prior over the CEG space

For any practical problem \mathcal{C} , the set of all possible CEGs for a given ET, is likely to be a very large set, making setting a value for $q(C), \forall C \in \mathcal{C}$ a non-trivial task. An obvious way to set a non-informative or exploratory prior is to choose the uniform prior, so that $q(C) = \frac{1}{|\mathcal{C}|}$. This has the advantages of being simple to set and of eliminating the $\log q(C_1) - \log q(C_2)$ term in Equation (8).

A more sophisticated approach is to consider which potential clusters are more or less likely a priori, according to structural or causal beliefs, and to exploit the modular nature of CEGs by stating that the prior log Bayes factor of a CEG relative to C_0 is the sum of the prior log Bayes factors of the individual clusters relative to their components completely unclustered, and that these priors are modular across CEGs. This approach makes it simple to elicit priors over \mathcal{C} from a lay expert, by requiring the elicitation only of the prior probability of each possible stage.

A particular computational benefit of this approach is when the prior Bayes factor of any CEG C with C_0 is believed to be zero, because one or more of its clusters is considered to be impossible. This is equivalent in the algorithm to not including the CEG in its search at all, as though it was never in \mathcal{C} in the first place, with the obvious simplification of the search following.

4.3. The prior over the parameter space

Just as when attempting to set $q(C)$, the size of most CEGs in practise leads to intractability of setting $p(\mathbf{x}|C)$ for each CEG C individually. However, the task is again made possible by exploiting the structure of a CEG with judicious modelling assumptions.

Assuming independence between the likelihoods of the stages for every CEG, so that $p(\mathbf{x}|\boldsymbol{\pi}, C)$ is as determined by Equation (1), and the fact that $p(\mathbf{x}|C) = \int p(\mathbf{x}|\boldsymbol{\pi}, C)p(\boldsymbol{\pi}|C)d\boldsymbol{\pi}$, it is clear that to set the marginal likelihood for each CEG is equivalent to setting the prior over the CEG's parameters, i.e. setting $p(\boldsymbol{\pi}|C)$ for each C . With the two further structural assumptions that the stage priors are independent for all CEGs (so that $p(\boldsymbol{\pi}|C) = \prod_{i=1}^k p(\boldsymbol{\pi}_i|C)$) and that equivalent stages in different CEGs have the same prior distributions on their probability vectors, (i.e. $p(\boldsymbol{\pi}_i|C_1) = p(\boldsymbol{\pi}_i|C_2)$), it can be seen that the problem of setting $p(\mathbf{x}|\boldsymbol{\pi}, C)$ is reduced to setting the parameter priors of each non-trivial floret in C_0 ($p(\boldsymbol{\pi}_i|C_0), i = 1, \dots, k$) and the parameter priors of stages that are clusters of stages of C_0 .

The usual prior put on the probability parameters of finite discrete BNs is the product Dirichlet distribution. In Geiger and Heckerman [14] the surprising result was shown that a product Dirichlet prior is inevitable if local and global independence are assumed to hold over all Markov equivalent graphs on at least two variables. In this paper we show that a similar characterisation can be made for CEGs given the assumptions in the previous paragraph. We will first show that the floret parameters in C_0 must have Dirichlet priors, and second that all CEGs formed by clustering the florets in C_0 have Dirichlet priors on the stage parameters. One characterisation of C_0 is given by Theorem 5.

Theorem 5. *If it is assumed a priori that the rates at which units take the root-to-leaf paths in C_0 are independent ("path independence") and that the probability of which edge units take after arriving at a situation v is independent of the rate at which units arrive at v ("floret independence"), then the non-trivial florets of C_0 have independent Dirichlet priors on their probability vectors.*

PROOF. The proof is in the Appendix.

Thus $p(\boldsymbol{\pi}_i|C_0)$ is entirely determined by the stated rates $\gamma(\lambda)$ on the root-to-leaf paths $\lambda \in \Lambda(C_0)$ of C_0 . This is similar to the "equivalent sample sizes" method of assessing prior uncertainty of Dirichlet hyperparameters in BNs as discussed in Section 2 of Heckerman [8].

Another way to show that all non-trivial situations in C_0 have Dirichlet priors on their parameter spaces is to use the characterisation of the Dirichlet distribution first proven by Geiger and Heckerman [14], repeated here as Theorem 6.

Theorem 6. *Let $\{\theta_{ij}\}, 1 \leq i \leq k, 1 \leq j \leq n, \sum_{ij} \theta_{ij} = 1$, where k and n are integers greater than 1, be positive random variables having a strictly positive pdf $f_U(\{\theta_{ij}\})$. Define $\theta_{i\cdot} = \sum_{j=1}^n \theta_{ij}$, $\theta_{\cdot i} = \{\theta_{ij}\}_{j=1}^{n-1}$, $\theta_{j|i} = \theta_{ij}/\sum_{j=1}^n \theta_{ij}$, and $\theta_{j|i} = \{\theta_{ij}\}_{j=1}^{n-1}$.*

Then if $\{\theta_{i\cdot}, \theta_{j|1}, \dots, \theta_{j|k}\}$ are mutually independent, $f_U(\{\theta_{ij}\})$ is Dirichlet.

PROOF. Theorem 2 of Geiger and Heckerman [14].

Corollary 7. *If C_0 has a composite number m of root-to-leaf paths and all Markov equivalent CEGs have independent floret distributions then the vector of probabilities on the root-to-leaf paths of C_0 must have a Dirichlet prior. This means in particular that, from the properties of the Dirichlet distribution, the floret of each situation with at least two outgoing edges has a Dirichlet prior on its edges.*

PROOF. Construct an event tree C'_0 with m root-to-leaf paths, where the floret of the root node v'_0 has k edges and each of the florets extending from the children of v'_0 have n edges terminating in leaf nodes, where $m = kn, k \geq 2, n \geq 2$. This will always be possible with a composite m . C'_0 describes the same atomic events as C_0 with a different decomposition.

Let the random variable associated with the root floret of C'_0 be X , and let the random variable associated with each of the other florets be $Y|X = i, i = 1, \dots, k$. Let $\theta_{ij} = P(X = i, Y = j)$. Then by the definition of event trees, $P(\theta_{ij} > 0) > 0, 1 \leq i \leq k, 1 \leq j \leq n$ and $\sum \theta_{ij} = 1$. By the notation of Theorem 6, $\theta_{i\cdot} = P(X = i)$ and $\theta_{\cdot j|i} = P(Y = j|X = i)$.

By hypothesis the floret distributions of C'_0 are independent. Therefore the condition of Theorem 6 holds and hence $f_U(\theta_{ij})$ is Dirichlet. From the equivalence of the atomic events, the probability distribution over the root-to-leaf path probabilities of C_0 is also Dirichlet, and so by Lemma 16, all non-trivial florets of C_0 therefore have Dirichlet priors on their probability vectors.

To show that the stage parameters of all the other CEGs in \mathbf{C} have independent Dirichlet priors, an inductive approach will be taken. Because of the assumption of consistency – that two identically composed stages in different CEGs have identical priors on their parameter space – for any given CEG C whose stages all have independent Dirichlet priors on their parameters spaces, it is known that another CEG C^* formed by clustering two stages u_{1c}, u_{2c} from C into one stage u_{c^*} will have independent Dirichlet priors on all its stages apart from u_{c^*} . It is thus only required to show that π_{c^*} has a Dirichlet prior. We prove this result for a class of CEGs called REGULAR CEGs.

Definition 8. A stage u is REGULAR if and only if every path $\lambda \in \Lambda(C)$ contains either one situation in u or none of the situations in u .

Definition 9. A CEG is REGULAR if and only if every situation $u \in \mathbf{u}(C)$ is regular.

Theorem 10. Let C be a regular CEG, and let C^* be the CEG that is formed from C by setting two of its stages, u_{1c} and u_{2c} , as being in the same stage u_{c^*} , where u_{c^*} is a regular stage, with all other attributes of the CEG unchanged from C .

If all stages in C have Dirichlet priors, then assuming that equivalent stages in different CEGs have equivalent priors, all stages in C^* have Dirichlet priors.

PROOF. Without loss of generality, let all situations in u_{1c} and u_{2c} have s children each, and let the total number of situations in u_{1c} and u_{2c} be r . Thus there are r situations in u_{c^*} , each with s children. By the assumption of prior consistency across stages, all stages in C^* have Dirichlet priors on their parameter spaces, so it is only required to prove that u_{c^*} has a Dirichlet prior.

Consider the CEG C' formed as follows: Let the root node of C' , v_0 , have 2 children, v_1 and v' . Let v' be a terminal node, and let v_1 have r children, $\{v_1(1), \dots, v_1(r)\}$, which are equivalent to the situations in u_{c^*} , including the property that they are in the same stage $u_{c'}$. Lastly, let the children of $\{v_1(1), \dots, v_1(r)\}$, $\{v_1(1, 1), \dots, v_1(1, s), \dots, v_1(r, 1), \dots, v_1(r, s)\}$, be leaf nodes in C' .

By construction, the prior for $u_{c'}$ is the same as that for u_{c^*} .

Now construct another CEG $C^{*'}$ from C' by reversing the order of the stages v_1 and $u_{c'}$. The new CEG has root node v_0 with the same distribution as $v_0 \in C'$. v_0 now has two children v' – the same as before – and v_2 , which has s children $\{v_2(1), \dots, v_2(s)\}$ in the same stage. Each node $v_2(i), i = 1, \dots, s$ has r children $v_2(i, 1), \dots, v_2(i, r)$, all of which are leaf nodes.

The two CEGs $C^{*'}$ and C' are Markov equivalent, as it is clear that $P(v_1(i, j)) = P(v_2(j, i)), i = 1, \dots, r, j = 1, \dots, s$. The probabilities on the floret of v_2 are thus equal to the probabilities of the situations in the stage of $u_{c'}$, and hence u_{c^*} . Because v_2 is a stage with only one situation, Theorem 5 implies that it has a Dirichlet prior. Therefore u_{c^*} has a Dirichlet prior.

An alternative justification for assigning a Dirichlet prior to any stage that is formed by clustering situations with Dirichlet priors on their state spaces can be obtained which does not depend on assuming Markov equivalency between CEGs derived from different event trees by assuming a property analogous to that of “parameter modularity” for BNs [15]. This property states that the distribution over structures common to two CEGs should be identical.

Definition 11. Let u be a stage in a CEG C composed of the situations v_1, \dots, v_n from C_0 , each of which has m children $v_{i1}, \dots, v_{im}, i = 1, \dots, n$ such that v_{ij} are the same colour for all i for each j . Then u has the property of MARGIN EQUIVALENCY if

$$\pi_{uj} = P(v_{1j} \text{ or } v_{2j} \text{ or } \dots \text{ or } v_{nj} | v_1 \text{ or } v_2 \text{ or } \dots \text{ or } v_n) \quad (9)$$

$$= \frac{\sum_{i=1}^n P(v_{ij})}{\sum_{i=1}^n P(v_i)} \quad (10)$$

is the same for both C and C_0 for $j = 1, \dots, m$.

Definition 12. C has margin equivalency if all of its stages have margin equivalency.

Theorem 13. Let u_c be a stage as defined in Definition 11 with $m \geq 2$. Then assuming independent priors between the situations for the associated finest-partition CEG C_0 of C , $\pi_{v_i} \sim \text{Dir}(\alpha_i)$ where $\alpha_i = (\alpha_{i1}, \dots, \alpha_{im})$ for each $v_i, i = 1, \dots, n$. Furthermore, for both C and C_0 , $\pi_u \sim \text{Dir}(\alpha_u)$, where $\alpha_u = (\sum_i \alpha_{i1}, \dots, \sum_i \alpha_{im})$.

PROOF. From Theorem [5] or Corollary [7], every non-trivial floret in C_0 has a Dirichlet prior on its edges, which includes in this case the situations v_1, \dots, v_n .

Let $\gamma_{ij} = \gamma \pi_{ij}$ for $i = 1, \dots, n, j = 1, \dots, m$ for some $\gamma \in \mathbb{R}^+$. Then it is a well-known fact that $\gamma_{ij} \sim \text{Gamma}(\alpha_{ij}, \beta)$ for all $1 \leq i \leq n, 1 \leq j \leq m$ for some $\beta > 0$, and that $\perp_j \gamma_{ij}$. As $\perp_i \pi_{v_i}, \perp_{ij} \gamma_{ij}$. Then by Lemma 15, letting $I[j]$ be the set of edges $\{e_{ij} = e(v_i, v_{ij}), i = 1, \dots, n\}$ for $j = 1, \dots, m$,

$$\pi_u \sim \text{Dir}(\sum_i \alpha_{i1}, \dots, \sum_i \alpha_{im})$$

By margin equivalency, π_u must be set the same way for C .

Note that the posterior of π_u for a stage u that is composed of the C_0 situations v_1, \dots, v_n is thus $\pi_u | \mathbf{x} \sim \text{Dir}(\alpha_u^*)$ where $\alpha_u^* = \alpha_u + \mathbf{x}_u = \sum_{i=1}^n \alpha_{v_n} + \sum_{i=1}^n \mathbf{x}_{v_n}$. Equation (8), therefore, becomes

$$\begin{aligned} \log \frac{q(C_1 | \mathbf{x})}{q(C_2 | \mathbf{x})} &= \log q(C_1) - \log q(C_2) + s(\alpha_{1a}) - s(\alpha_{1a}^*) + t(\alpha_{1a}^*) - t(\alpha_{1a}) \\ &\quad + s(\alpha_{1b}) - s(\alpha_{1b}^*) + t(\alpha_{1b}^*) - t(\alpha_{1b}) - s(\alpha_{1a} + \alpha_{1b}) \\ &\quad + s(\alpha_{1a}^* + \alpha_{1b}^*) - t(\alpha_{1a}^* + \alpha_{1b}^*) + t(\alpha_{1a} + \alpha_{1b}) \quad (11) \end{aligned}$$

4.4. The algorithm

The algorithm thus proceeds as follows:

1. Starting with the initial ET model, form the CEG C_0 with the finest possible partition, where all leaf nodes are placed in the terminal stage u_∞ and all nodes with only one emanating edge are placed in the same stage. Calculate $\log q(C_0|\mathbf{x})$ using (4).
2. For each pair of situations $v_i, v_j \in C_0$ with the same number of edges, calculate $\log \frac{q(C_1^*|\mathbf{x})}{q(C_0|\mathbf{x})}$ where C_1^* is the CEG formed by having v_i, v_j in the same stage and keeping all others in their own stage; do not calculate if $q(C_1^*) = 0$.
3. Let $C_1 = \max_{C_1^*} (\log \frac{q(C_1^*|\mathbf{x})}{q(C_0|\mathbf{x})})$.
4. Now calculate C_2^* for each pair of stages in C_1 except where $q(C_2^*) = 0$, and record $C_2 = \max(q(C_2^*|\mathbf{x}))$.
5. Continue for C_3, C_4 and so on until the coarsest partition C_∞ has been reached.
6. Find $C = \max(C_0, C_1, \dots, C_\infty)$, and select this as the MAP model.

We note that the algorithm can also be run backwards, starting from C_∞ and splitting one cluster in two at each step. This has the advantage of making the identification of positions in the MAP model easier.

5. Examples

5.1. Simulated data

To first demonstrate the efficacy of the algorithm described above we implement the algorithm using simulated data for Example 1, where the CEG generating the data was as known and described in Section 1. Figure 5.1 shows the number of students in the sample who reached each situation in the tree.

In this complete dataset the progress of 1000 students has been tracked through the event tree. Half are assigned to take module *A* first and the other half *B*. By finding the MAP CEG model in the light of this data we may find out whether the three hypotheses posed in the introduction are valid. We repeat them here for convenience:

1. The chances of doing well in the second component are the same whether the student passed first time or after a resit.
2. The components *A* and *B* are equally hard.
3. The distribution of marks for the second component is unaffected by whether students passed or got a distinction for the first component.

We set a uniform prior on the CEG priors and on the root-to-leaf paths of C_0 , the finest partition of the tree, for illustration purposes. The algorithm is then implemented as follows.

There are only two florets with two edges; with Beta(1,3) priors on each and a Beta(2,6) prior on the combined stage, the log Bayes factor is -1.85. Carrying out similar calculations for all the pairs of nodes with three edges, it is first decided to merge the nodes $P_{1,A}$ and $P_{1,B}$, which has a log Bayes factor of -3.76 against leaving them apart. Applying the algorithm to the updated set of nodes and iterating, the CEG in Figure 5.1 is found to be the MAP one.

Under this model, it can be seen that all three hypotheses above are satisfied and that the MAP model is the correct one.




Figure 4: The event tree from Example 1 with the numbers representing the number of students in a simulated sample who reached each situation.




Figure 5: The MAP CEG for that event tree in Figure 5.1

5.2. Student test data

In our second example we apply the learning algorithm to a real dataset in order to test the algorithm’s efficacy in a real-life situation and to identify remaining issues with its usage. The dataset we used was an appropriately disguised set of marks taken over a 10-year period from four core modules of the MORSE degree course taught at the University of Warwick. A part of the event tree used as the underlying model for the first two modules is shown in Figure 5.2, along with a few illustrative data points. This is a simplification of a much larger study that we are currently investigating but large enough to illustrate the richness of inference possible with our model search.

For simplicity, the prior distributions on the candidate models and on the root-to-leaf paths for C_0 were both chosen to be uniform distributions.

The MAP CEG model was not C_0 , so that there were some non-trivial stages. In total, 170 situations were clustered into 32 stages. Some of the more interesting stages of this model are described in Table 1.

Stage	Probability vector	Students	Situations	Locations	Comments
7	(0.47, 0.44, 0.08)	685	2	1; 1,1,1	High achievers
11	(0.22, 0.43, 0.35)	412	6	3; 1,2; 3,1; 1,1,3	Middling students
13	(0.33, 0.33, 0.33)	16	18	4; 4,2; 4,3	No students appeared in 17 of these situations
17	(0.07, 0.27, 0.66)	86	4	1,3; 3,2; 3,2,4	Struggling students
27	(0.19, 0.56, 0.25)	464	7	1,1,4; 1,2,2; 1,3,2; 1,4,2	More likely to get grade 2 than stage 11
28	(0.11, 0.51, 0.38)	436	6	1,2,3; 3,1,3; 1,2,4	More likely to get grade 3 than stage 27

Table 1: Selected stages of MAP CEG model formed from data described in Section 5.2. The columns respectively detail the stage number, posterior expectation of the probability vector of that stage (rounded to two decimal places), number of students passing through that stage in the dataset, number of situations from the original ET in that stage, examples of situations in that stage (shown as sequence of grades, where “4” means that grade is missing), and any comments or observations related to that stage.




Figure 6: Sub-tree of the event tree of possible grades for the MORSE degree course at the University of Warwick. Each floret of two edges describes whether a student's marks are available for a particular module (denoted by the edge labelled A for the first module) or whether they are missing (NA). If they are available, then they are counted as grade 1 if are 70% or higher, grade 2 if they are between 50% and 69% inclusive, and grade 3 if they are below 50%. Some illustrative count data are shown on corresponding nodes.

From inspecting the membership of stages it was possible to identify various situations which were discovered to share distributions. From example, students who reach one of the two situations in stage 7 have an expected probability of 0.47 in getting a high mark, an expected probability of 0.44 of getting a middling grade, and only an expected probability of 0.08 of achieving the lowest grade. From being in a stage of their own, it can be deduced that students in these situations have qualitatively different prospects from students in any other situations. In contrast, students who reach one of the four situations in stage 17 have an expected probability of 0.66 of getting the lowest grade.

6. Discussion

In this paper we have shown that chain event graphs are not just an efficient way of storing the information contained in an event tree, but also a natural way to represent the information that is most easily elicited from a domain expert: the order in which events happen, the distributions of variables conditional on the process up to the point they are reached, and prior beliefs about the relative homogeneity of different situations. This strength is exploited when the MAP CEG is discovered, as this can be used in a qualitative fashion to detect homogeneity between seemingly disparate situations.

There are a number extensions to the theory in this paper that are currently being pursued. These fall mostly into the two categories: creating even richer model classes than those considered here; and developing even more efficient algorithms for selecting the MAP model in these model classes.

The first category includes dynamic chain event graphs. This framework can supply a number of different model classes. The simplest case involves selecting a CEG structure that is constant across time, but with a time series on its parameters. A bigger class would allow the MAP CEG structure to change over time. These larger model classes would clearly be useful in the educational setting considered in this paper, as they would allow for background changes in the students' abilities, for example.

Another important model class is that which arises from uncertainty about the underlying event tree. A similar model search algorithm to the one described in this paper is possible in this case after setting a prior distribution on the candidate event trees.

In order to search any of these model classes more effectively, the problem of finding the MAP model can be reformulated as a weighted MAX-SAT problem, for which algorithms have been developed. This approach was used to great effect for finding a MAP BN by Cussens [16].

Appendix

Theorem 5 is based on three well-known results concerning properties of the Dirichlet distribution, which we review below.

Lemma 14. *Let $\gamma_j \sim \text{Gamma}(\alpha_j, \beta)$, $j = 1, \dots, n$ where $\alpha_j > 0$ for $j \in \{1, \dots, n\}$, $\beta > 0$ and $\prod_{i \in \{1, \dots, n\}} \gamma_i$. Furthermore, let $\theta_j = \frac{\gamma_j}{\gamma}$ for $j \in \{1, \dots, n\}$, where $\gamma = \sum_{i=1}^n \gamma_i$. Then $\theta = (\theta_i)_{i \in \{1, \dots, n\}} \sim \text{Dir}(\alpha_1, \dots, \alpha_n)$.*

PROOF. Kotz et al [17].

Lemma 15. *Let $I[j] \subseteq \{1, \dots, n\}$, $\gamma(I[j]) = \sum_{i \in I[j]} \gamma_i$ and $\theta(I[j]) = \sum_{i \in I[j]} \theta_i$. Then for any partition $I = \{I[1], \dots, I[k]\}$ of $\{1, \dots, n\}$,*

$$\theta(I) = (\theta(I[1]), \theta(I[2]), \dots, \theta(I[k])) \sim \text{Dir}(\alpha(I[1]), \dots, \alpha(I[k]))$$

where $\alpha(I[j]) = \sum_{i \in I[j]} \alpha_i$.

PROOF. For any $I[j] \subseteq \{1, \dots, n\}$, $\prod_{i \in I[j]} \gamma_i$, $\gamma(I[j]) \sim \text{Gamma}(\alpha(I[j]), \beta)$ (a well-known result; see, for example, Weatherburn [18]), and for any partition $I = \{I[1], \dots, I[k]\}$ of $\{1, \dots, n\}$, $\prod_{i \in \{1, \dots, k\}} \gamma(I[j])$. Therefore, as

$$\theta(I[j]) = \sum_{i \in I[j]} \theta_i = \sum_{i \in I[j]} \frac{\gamma_i}{\gamma} = \frac{\gamma(I[j])}{\gamma}, \quad j = 1, \dots, k$$

and $\gamma = \sum_{i=1}^k \gamma(I[i])$, the result follows from Lemma 14.

Lemma 16. *For any $I[j] \subseteq \{1, \dots, n\}$ where $|I[j]| \geq 2$,*

$$\theta_{I[j]} = \left(\frac{\theta_i}{\theta(I[j])} \right)_{i \in I[j]} \sim \text{Dir}((\alpha_i)_{i \in I[j]})$$

PROOF. Wilks [19].

Theorem 17. *Let the rates of units along the root-to-leaf paths $\lambda_i \in \Lambda$, $i \in \{1, \dots, |\Lambda|\}$ of an event tree T have independent Gamma distributions with the same scale parameter, i.e. $\gamma_i = \gamma(\lambda_i) \sim \text{Gamma}(\alpha_i, \beta)$, $i \in \{1, \dots, |\Lambda|\}$ and $\prod_{i \in \{1, \dots, |\Lambda|\}} \gamma_i$. Then the distribution on each floret in the tree will be Dirichlet.*

PROOF. Consider a floret \mathcal{F} with root node v and edge set $\{e_1, \dots, e_l\}$. The rate for each edge e_i , $\gamma(e_i)$, is equal to $\gamma(\lambda_{e_i})$, where λ_{e_i} is the root-to-leaf path that intersects with e_i , so that $\gamma(e_i) \sim \text{Gamma}(\alpha_{e_i}, \beta)$ and $\prod_{i \in \{1, \dots, l\}} \gamma(e_i)$.

Let $I = \{I[\mathcal{F}], I[\overline{\mathcal{F}}]\}$ partition Λ , where $I[\mathcal{F}] = \{\lambda_{e_1}, \dots, \lambda_{e_l}\}$ and $I[\overline{\mathcal{F}}] = I - I[\mathcal{F}]$. Then by Lemma 16, the probability vector on \mathcal{F} is Dirichlet, where

$$\theta_{I[\mathcal{F}]} \sim \text{Dir}((\alpha_{e_i})_{i \in \{1, \dots, l\}})$$

References

- [1] R. G. Cowell, A. P. Dawid, S. L. Lauritzen, D. J. Spiegelhalter, Probabilistic Networks and Expert Systems, Springer, 1999.
- [2] S. L. Lauritzen, Graphical Models (Oxford Statistical Science Series), Oxford University Press, USA, 1996.
- [3] D. Poole, N. L. Zhang, Exploiting contextual independence in probabilistic inference, J. Artificial Intelligence Res. 18 (2003) 263–313.

- [4] J. Q. Smith, P. E. Anderson, Conditional independence and chain event graphs, *Artificial Intelligence* 172 (1) (2008) 42–68.
- [5] G. Shafer, *The Art of Causal Conjecture*, Artificial Intelligence, The MIT Press, 1996.
- [6] D. G. T. Denison, C. C. Holmes, B. K. Mallick, A. F. M. Smith, *Bayesian Methods for Nonlinear Classification and Regression*, Wiley Series in Probability and Statistics, Wiley, 2002.
- [7] R. Castelo, The discrete acyclic digraph markov model in data mining, Ph.D. thesis, Faculteit Wiskunde en Informatica, Universiteit Utrecht (Apr. 2002).
- [8] D. Heckerman, A tutorial on learning with bayesian networks, in: M. I. Jordan (Ed.), *Learning in Graphical Models*, MIT Press, 1999, pp. 301–354.
- [9] D. J. Spiegelhalter, S. L. Lauritzen, Sequential updating of conditional probabilities on directed graphical structures, *Networks* 20 (5) (1990) 579–605.
- [10] J. Bernardo, A. F. M. Smith, *Bayesian Theory*, Wiley, Chichester, England, 1994.
- [11] J. W. Lau, P. J. Green, Bayesian Model-Based clustering procedures, *Journal of Computational and Graphical Statistics* 16 (3) (2007) 526–558.
- [12] S. Richardson, P. J. Green, On bayesian analysis of mixtures with an unknown number of components, *Journal of the Royal Statistical Society. Series B (Methodological)* 59 (4) (1997) 731–792.
- [13] N. A. Heard, C. C. Holmes, D. A. Stephens, A quantitative study of gene regulation involved in the immune response of anopheline mosquitoes: An application of bayesian hierarchical clustering of curves, *Journal of the American Statistical Association* 101 (473) (2006) 18–29.
- [14] D. Geiger, D. Heckerman, A characterization of the dirichlet distribution through global and local parameter independence, *The Annals of Statistics* 25 (3) (1997) 1344–1369.
- [15] D. Heckerman, M. P. Wellman, Bayesian networks, *Communications of the ACM* 38 (3) (1995) 27–30.
- [16] J. Cussens, Bayesian network learning by compiling to weighted MAX-SAT, in: D. A. McAllester, P. Myllymäki (Eds.), *Proceedings of the 24th Conference in Uncertainty in Artificial Intelligence*, AUAI Press, Helsinki, Finland, 2008, pp. 105–112.
- [17] S. Kotz, N. Balakrishnan, N. L. Johnson, *Continuous Multivariate Distributions*, 2nd Edition, Wiley series in probability and statistics. Applied probability and statistics, Wiley, New York, 2000.
- [18] C. E. Weatherburn, *A first course in mathematical statistics*, 2nd Edition, CUP, 1949.
- [19] S. S. Wilks, *Mathematical Statistics*, Wiley, New York, 1962.